

- netty
 - 背景
 - 问题剖析
 - netty 流程剖析
 - 连接建立
 - 接受请求
 - 业务处理
 - 响应体写入
 - 总体逻辑
 - 内存剖析
 - 内存核心机制
 - 内存分类与池化技术
 - 内存分配与回收机制
 - 零拷贝优化
 - 背景问题解决方案

netty

背景

场景：业务线程生成数据速度远高于IO线程发送速度，ChannelOutboundBuffer队列积压。

原理：writeAndFlush操作将数据封装为Task加入队列，若队列无限制增长，触发堆内存OOM。

问题剖析

netty 流程剖析

服务端请求流程图请求过程

建立连接 ---> 接受请求 ---> 处理请求 ---> 回写响应体 ---> 关闭连接

连接建立

1. channel注册到selector上

```
AbstractNioChannel.java

@Override
protected void doRegister() throws Exception {
    boolean selected = false;
    for (;;) {
        try {
            //把channel注册到selector上
            selectionKey = javaChannel().register(eventLoop().unwrappedSelector(), 0, this);
```

```

        return;
    } catch (CancelledKeyException e) {
        if (!selected) {
            // Force the Selector to select now as the "canceled" SelectionKey may still
            // be cached and not removed because no Select.select(..) operation was called
            // yet.
            eventLoop().selectNow();
            selected = true;
        } else {
            // We forced a select operation on the selector before but the SelectionKey is
            // still cached
            // for whatever reason. JDK bug ?
            throw e;
        }
    }
}
}
}

```

void register(EventLoop eventLoop, final ChannelPromise promise) //eventloop() 在注册的时候注册的是eventLoop

注册到BOSS线程中

```

final ChannelFuture initAndRegister() {
    Channel channel = null;
    try {
        //创建一个channel对象
        channel = channelFactory.newChannel();
        //初始化channel配置
        init(channel);
    } catch (Throwable t) {
        if (channel != null) {
            //如果 newChannel 崩溃, channel 可以为 null
            channel.unsafe().closeForcibly();
            // 由于通道尚未注册, 我们需要强制使用 GlobalEventExecutor
            return new DefaultChannelPromise(channel,
GlobalEventExecutor.INSTANCE).setFailure(t);
        }
        // 由于通道尚未注册, 我们需要强制使用 GlobalEventExecutor
        return new DefaultChannelPromise(new FailedChannel(),
GlobalEventExecutor.INSTANCE).setFailure(t);
    }
    //获取初始化设置的bossGroup,将channel绑定到
    ChannelFuture regFuture = config().group().register(channel);
    if (regFuture.cause() != null) {
        if (channel.isRegistered()) {
            channel.close();
        } else {
            channel.unsafe().closeForcibly();
        }
    }
}
/**
 * 如果在这promise没有失败, 则一定是以下原因
 * 1.如果我们尝试从事件循环中注册, 此时注册已经完成。因为channel注册完成, 使用bind(),connect () 是
安全的
 * 2.如果我们尝试从其他线程注册, 则注册请求已经成功添加到事件循环的任务队列以供稍后执行
 * 因为bind(),connect () 将在定时任务后执行
 * 应为register(), bind(), and connect()被绑定在相同的线程
 */
return regFuture;
}

```

2. 事件监听激活在Channel绑定端口后, 通过ChannelPipeline触发channelActive事件, 最终调用AbstractNioChannel#doBeginRead设置OP_ACCEPT监听:

触发事件后调用ProcessSelectedKey, Reactor的Selector 内核态的tcp accept 转为用户态的Event事件

```
private void processSelectedKey(SelectionKey k, AbstractNioChannel ch) {
    final NioUnsafe unsafe = ch.unsafe();
    int readyOps = k.readyOps();
    if ((readyOps & (SelectionKey.OP_READ | SelectionKey.OP_ACCEPT)) != 0) {
        unsafe.read(); // 核心处理入口
    }
}
```

NioMessageUnsafe#read() (NioServerSocketChannel的Unsafe实现) 负责处理OP_ACCEPT事件:

```
// NioMessageUnsafe#read
public void read() {
    try {
        do {
            SocketChannel socketChannel = javaChannel().accept();
            // 创建客户端Channel并触发pipeline传播
            pipeline.fireChannelRead(socketChannel);
        } while (继续接受新连接);
    } catch (IOException e) { /* 异常处理 */ }
}
```

ServerSocketChannel.accept()接受客户端连接, 生成SocketChannel。

新连接的SocketChannel被封装为NioSocketChannel, 并通过pipeline.fireChannelRead()传播到业务逻辑处理器。

新连接注册到Worker线程

在ChannelRead事件传播中, 通过ServerBootstrap的ServerBootstrapAcceptor将新NioSocketChannel注册到Worker线程组的NioEventLoop, 监听OP_READ事件:

```
// ServerBootstrapAcceptor#channelRead
@Override
@SuppressWarnings("unchecked")
public void channelRead(ChannelHandlerContext ctx, Object msg) {
    //对应传入进来的SocketChannel(创建连接的时候创建的)
    final Channel child = (Channel) msg;
    //最外侧添加的自己的写的handler
    child.pipeline().addLast(childHandler);
    //设置属性
    setChannelOptions(child, childOptions, logger);
    setAttributes(child, childAttrs);

    try {
        //对应WorkGroup
        childGroup.register(child).addListener(new ChannelFutureListener() {
            @Override
            public void operationComplete(ChannelFuture future) throws Exception {
                if (!future.isSuccess()) {
                    forceClose(child, future.cause());
                }
            }
        });
    } catch (Throwable t) {
        forceClose(child, t);
    }
}

private static void forceClose(Channel child, Throwable t) {
    child.unsafe().closeForcibly();
    logger.warn("Failed to register an accepted channel: {}", child, t);
}
```

1. 注册到Worker线程

NioEventLoop#register

```
public void register(final SelectableChannel ch, final int interestOps, final NioTask<?>
task) {
    ObjectUtil.checkNotNull(ch, "ch");
    if (interestOps == 0) {
        throw new IllegalArgumentException("interestOps must be non-zero.");
    }
    if ((interestOps & ~ch.validOps()) != 0) {
        throw new IllegalArgumentException(
            "invalid interestOps: " + interestOps + "(validOps: " + ch.validOps() + ')');
    }
    ObjectUtil.checkNotNull(task, "task");

    if (isShutdown()) {
        throw new IllegalStateException("event loop shut down");
    }

    if (inEventLoop()) {
        register0(ch, interestOps, task);
    } else {
        try {
            // Offload to the EventLoop as otherwise
            java.nio.channels.spi.AbstractSelectableChannel.register
            // may block for a long time while trying to obtain an internal lock that may be
            hold while selecting.
            submit(new Runnable() {
                @Override
                public void run() {
                    register0(ch, interestOps, task);
                }
            }).sync();
        } catch (InterruptedException ignore) {
            // Even if interrupted we did schedule it so just mark the Thread as interrupted.
            Thread.currentThread().interrupt();
        }
    }
}
```

2. 监听事件

```
private void processSelectedKey(SelectionKey k, AbstractNioChannel ch) {
    final AbstractNioChannel.NioUnsafe unsafe = ch.unsafe();
    if (!k.isValid()) {
        final EventLoop eventLoop;
        try {
            eventLoop = ch.eventLoop();
        } catch (Throwable ignored) {
            // If the channel implementation throws an exception because there is no event
            loop, we ignore this
            // because we are only trying to determine if ch is registered to this event loop
            and thus has authority
            // to close ch.
            return;
        }
        // Only close ch if ch is still registered to this EventLoop. ch could have
        deregistered from the event loop
        // and thus the SelectionKey could be cancelled as part of the deregistration process,
        but the channel is
        // still healthy and should not be closed.
        // See https://github.com/netty/netty/issues/5125
        if (eventLoop == this) {
            // close the channel if the key is not valid anymore
            unsafe.close(unsafe.voidPromise());
        }
    }
}
```

```

        return;
    }

    try {
        int readyOps = k.readyOps();
        // We first need to call finishConnect() before try to trigger a read(...) or
        write(...) as otherwise
        // the NIO JDK channel implementation may throw a NotYetConnectedException.
        if ((readyOps & SelectionKey.OP_CONNECT) != 0) {
            // remove OP_CONNECT as otherwise Selector.select(..) will always return without
            blocking
            // See https://github.com/netty/netty/issues/924
            int ops = k.interestOps();
            ops &= ~SelectionKey.OP_CONNECT;
            k.interestOps(ops);

            unsafe.finishConnect();
        }

        // Process OP_WRITE first as we may be able to write some queued buffers and so free
        memory.
        <!-- 数据过多无法一次性写入的时候会由OP_WRITE事件产生 -->
        if ((readyOps & SelectionKey.OP_WRITE) != 0) {
            // Call forceFlush which will also take care of clear the OP_WRITE once there is
            nothing left to write
            ch.unsafe().forceFlush();
        }

        // Also check for readOps of 0 to workaround possible JDK bug which may otherwise lead
        // to a spin loop
        //请求读处理（断开连接）或接入连接
        //2种请求方式对应NioMessageChannel和NioByteChannel 2种对应的处理逻辑
        if ((readyOps & (SelectionKey.OP_READ | SelectionKey.OP_ACCEPT)) != 0 || readyOps ==
        0) {
            unsafe.read();
        }
    } catch (CancelledKeyException ignored) {
        unsafe.close(unsafe.voidPromise());
    }
}

```

3. unsafe.read

unsafe是NioByteUnsafe实例（定义在AbstractNioByteChannel中），负责底层SocketChannel的I/O操作。

```

public final void read() {
    final ChannelConfig config = config();
    if (shouldBreakReadReady(config)) {
        clearReadPending();
        return;
    }
    final ChannelPipeline pipeline = pipeline();
    //ByteBuf分配器
    final ByteBufAllocator allocator = config.getAllocator();
    //抉择下一次分配多少ByteBuf
    final RecvByteBufAllocator.Handle allocHandle = recvBufAllocHandle();
    allocHandle.reset(config);

    ByteBuf byteBuf = null;
    boolean close = false;
    try {
        do {
            //分配适合的大小 ---- 内存分配策略核心部分 ---- AdaptiveRecvByteBufAllocator
            byteBuf = allocHandle.allocate(allocator);
            //读并且记录读了多少，如果堵满了，下次continue的化就直接扩容
            allocHandle.lastBytesRead(doReadBytes(byteBuf));
            if (allocHandle.lastBytesRead() <= 0) {
                // nothing was read. release the buffer.
                byteBuf.release();
            }
        }
    }
}

```

```

        byteBuf = null;
        close = allocHandle.lastBytesRead() < 0;
        if (close) {
            // There is nothing left to read as we received an EOF.
            readPending = false;
        }
        break;
    }

    allocHandle.incMessagesRead(1);
    readPending = false;
    //pipeline上执行，业务逻辑的处理就在这个地方
    pipeline.fireChannelRead(byteBuf);
    byteBuf = null;
} while (allocHandle.continueReading());
//计算这次事件总共读了多少数据，计算下次分配大小
allocHandle.readComplete();
//相当于完成本次读事件处理
pipeline.fireChannelReadComplete();

if (close) {
    closeOnRead(pipeline);
}
} catch (Throwable t) {
    handleReadException(pipeline, byteBuf, t, close, allocHandle);
} finally {
    // Check if there is a readPending which was not processed yet.
    // This could be for two reasons:
    // * The user called Channel.read() or ChannelHandlerContext.read() in
channelRead(...) method
    // * The user called Channel.read() or ChannelHandlerContext.read() in
channelReadComplete(...) method
    //
    // See https://github.com/netty/netty/issues/2254
    if (!readPending && !config.isAutoRead()) {
        removeReadOp();
    }
}
}
}

```

业务处理

DefaultChannelPipeline 添加各类Handler后在不同阶段会触发不同的方法 以下是Netty中 **DefaultChannelPipeline**在不同阶段触发的事件方法分类及对应场景的整理表格，结合入站（Inbound）和出站（Outbound）事件类型进行说明：

阶段/场景	事件类型	触发方法	传播方向	说明
通道初始化	Inbound	fireChannelRegistered()	从 Head 到 Tail	通道注册到EventLoop后触发，用于初始化资源（如 ChannelInitializer ）
	Inbound	fireChannelActive()	从 Head	通道激活（如TCP连接建立或端口绑定完成）时触发

阶段/场景	事件类型	触发方法	传播方向	说明
			到Tail	
数据接收 (读操作)	Inbound	<code>fireChannelRead(Object msg)</code>	从Head到Tail	读取到数据时触发，需在Handler中调用 <code>ctx.fireChannelRead()</code> 传递
	Inbound	<code>fireChannelReadComplete()</code>	从Head到Tail	单次读取循环完成时触发，用于批量处理后的逻辑（如刷新响应）
数据发送 (写操作)	Outbound	<code>bind(SocketAddress) / connect(SocketAddress)</code>	从Tail到Head	绑定端口或发起连接时触发，最终由HeadContext执行系统调用
	Outbound	<code>write(Object msg) / writeAndFlush(Object msg)</code>	从Tail到Head	数据写入缓冲区， <code>flush()</code> 触发底层发送
连接状态变更	Inbound	<code>fireChannelInactive()</code>	从Head到Tail	连接关闭时触发（如TCP连接断开）
	Inbound	<code>fireChannelUnregistered()</code>	从Head到Tail	通道从EventLoop注销时触发，用于资源清理
异常处理	Inbound	<code>fireExceptionCaught(Throwable cause)</code>	从Head到Tail	处理过程中发生异常时触发，默认由TailContext记录日志
流量控制	Inbound	<code>fireChannelWritabilityChanged()</code>	从Head到Tail	发送缓冲区水位变化时触发（如高水位线阻塞）
用户自定	Inbound	<code>fireUserEventTriggered(Object event)</code>	从Head	用户通过 <code>ctx.fireUserEventTriggered()</code>

阶段/场景	事件类型	触发方法	传播方向	说明
入站事件			到Tail	触发的自定义事件
出站操作控制	Outbound	<code>read() / deregister() / close()</code>	从Tail到Head	主动发起读取、注销或关闭连接操作

补充说明：

- 1. 传播方向：
 - 入站事件（如数据接收）从HeadContext向TailContext传播。
 - 出站事件（如数据发送）从当前Handler向HeadContext传播，最终由Unsafe类执行底层I/O操作。
- 2. 关键源码逻辑：
 - 事件触发入口：NioEventLoop通过processSelectedKey()检测I/O事件，调用pipeline.fireXXX()方法。
 - 动态处理器链：通过addLast()将Handler封装为AbstractChannelHandlerContext节点，插入双向链表。
- 3. 设计亮点：
 - 责任链模式：通过双向链表实现事件动态编排，支持业务逻辑分层处理（如编解码、业务逻辑、日志）。
 - 线程安全：通过synchronized块保证多线程下链表操作安全。

如需深入源码细节，可参考DefaultChannelPipeline类中的fireChannelActive()和bind()方法实现。

响应体写入

通过ChannelHandlerContext或Channel调用writeAndFlush(), 本质上触发AbstractChannelHandlerContext中的统一入口方法。无论从TailContext还是当前Handler调用，最终会通过findContextOutbound()向前查找出站处理器链。

触发TailContext.writeAndFlush()

- 事件传播至Encoder.encode()转换Java对象为ByteBuf
- HeadContext调用Unsafe.write()写入ChannelOutboundBuffer
- Entry加入链表，更新totalPendingSize及水位状态
- 调用flush()后，遍历Entry链表发送数据至Socket
- 移除已发送Entry，触发Promise回调

@Override
public ChannelFuture writeAndFlush(Object msg, ChannelPromise promise) {
 write(msg, true, promise);
 return promise;
}

private void write(Object msg, boolean flush, ChannelPromise promise) {
 ObjectUtil.checkNotNull(msg, "msg");
 try {
 if (isValidPromise(promise, true)) {
 ReferenceCountUtil.release(msg);
 // cancelled


```

        return;
    }
} catch (RuntimeException e) {
    ReferenceCountUtil.release(msg);
    throw e;
}

final AbstractChannelHandlerContext next = findContextOutbound(flush ?
    (MASK_WRITE | MASK_FLUSH) : MASK_WRITE);
//引用计数用，用来检测内存泄漏
final Object m = pipeline.touch(msg, next);
EventExecutor executor = next.executor();
if (executor.inEventLoop()) {
    if (flush) {
        next.invokeWriteAndFlush(m, promise);
    } else {
        next.invokeWrite(m, promise);
    }
} else {
    final WriteTask task = WriteTask.newInstance(next, m, promise, flush);
    if (!safeExecute(executor, task, promise, m, !flush)) {
        // We failed to submit the WriteTask. We need to cancel it so we decrement the
pending bytes
        // and put it back in the Recycler for re-use later.
        //
        // See https://github.com/netty/netty/issues/8343.
        task.cancel();
    }
}
}

@Override
public final void write(Object msg, ChannelPromise promise) {
    assertEventLoop();
    //判断channel是否已经关闭了
    ChannelOutboundBuffer outboundBuffer = this.outboundBuffer;
    if (outboundBuffer == null) {
        try {
            // release message now to prevent resource-leak
            ReferenceCountUtil.release(msg);
        } finally {
            // If the outboundBuffer is null we know the channel was closed and so
            // need to fail the future right away. If it is not null the handling of the rest
            // will be done in flush0()
            // See https://github.com/netty/netty/issues/2362
            safeSetFailure(promise,
                newClosedChannelException(initialCloseCause, "write(Object,
ChannelPromise)"));
        }
        return;
    }
}

int size;
try {
    // 零拷贝
    msg = filterOutboundMessage(msg);
    size = pipeline.estimatorHandle().size(msg);
    if (size < 0) {
        size = 0;
    }
} catch (Throwable t) {
    try {
        ReferenceCountUtil.release(msg);
    } finally {
        safeSetFailure(promise, t);
    }
    return;
}
}

```

```

//消息发送到buf里面
outboundBuffer.addMessage(msg, size, promise);
}

@Override
protected final Object filterOutboundMessage(Object msg) {
    if (msg instanceof ByteBuf) {
        ByteBuf buf = (ByteBuf) msg;
        if (buf.isDirect()) {
            return msg;
        }

        return newDirectBuffer(buf);
    }

    if (msg instanceof FileRegion) {
        return msg;
    }

    throw new UnsupportedOperationException(
        "unsupported message type: " + StringUtil.simpleClassName(msg) + EXPECTED_TYPES);
}

public void addMessage(Object msg, int size, ChannelPromise promise) {
    Entry entry = Entry.newInstance(msg, size, total(msg), promise);
    if (tailEntry == null) {
        flushedEntry = null;
    } else {
        Entry tail = tailEntry;
        tail.next = entry;
    }
    tailEntry = entry; // 添加到队尾
    if (unflushedEntry == null) {
        unflushedEntry = entry;
    }

    // increment pending bytes after adding message to the unflushed arrays.
    // See https://github.com/netty/netty/issues/1619
    incrementPendingOutboundBytes(entry.pendingSize, false);
}

public void addFlush() {
    // There is no need to process all entries if there was already a flush before and no new
    messages
    // where added in the meantime.
    //
    // See https://github.com/netty/netty/issues/2577
    Entry entry = unflushedEntry;
    if (entry != null) {
        if (flushedEntry == null) {
            // there is no flushedEntry yet, so start with the entry
            flushedEntry = entry;
        }
        do {
            flushed ++;
            if (!entry.promise.setUncancellable()) {
                // Was cancelled so make sure we free up memory and notify about the freed
                bytes
                int pending = entry.cancel();
                decrementPendingOutboundBytes(pending, false, true);
            }
            entry = entry.next;
        } while (entry != null);

        // All flushed so reset unflushedEntry
        unflushedEntry = null;
    }
}

```

```
    }  
}
```

总体逻辑

将上述流程总结如下述总体逻辑：

请求处理流程

1. EventLoop线程模型
- Boss-Worker分工：
 - Boss线程组：仅处理OP_ACCEPT事件，轻量化设计。
 - Worker线程组：处理OP_READ/OP_WRITE事件，绑定客户端Channel生命周期。
 - 线程选择策略：通过EventExecutorChooser（如轮询或幂等选择器）分配Channel到不同线程。
2. ChannelPipeline与Handler
- 动态编排：服务端通过ServerBootstrapAcceptor将客户端Channel的Pipeline初始化为用户配置的childHandler链。
 - 零拷贝优化：通过ByteBuf和FileRegion实现高效数据传输，减少内存复制。
1. 异步与Future机制ChannelFuture：连接操作返回ChannelFuture，支持addListener()异步回调

内存剖析

请求内存申请过程

内存申请

阶段	内存操作	释放时机	引用机制
OP_ACCEPT	分配元数据 (Channel/Pipeline)	Channel关闭时自动释放	框架管理
OP_READ	分配接收缓冲区 (PooledByteBuf)	默认由TailContext释放，或业务手动释放	引用计数
业务处理	可能分配新ByteBuf（如响应数据）	业务调用writeAndFlush()后由框架释放	引用计数+发送完成回调
OP_WRITE	转换堆内存为堆外内存（零拷贝优化）	数据发送完成后自动释放	框架自动管理
异常/关闭	释放所有未回收的ByteBuf	channelInactive()或exceptionCaught()	强制释放

内存核心机制

Netty 的内存管理机制是其高性能网络通信的核心，通过池化技术、零拷贝、引用计数等策略优化内存使用效率并降低GC压力。以下是其核心机制及实现原理的详细解析：

内存分类与池化技术

1. 堆内存（Heap Buffer）与直接内存（Direct Buffer）

- 堆内存：分配在JVM堆上，受GC管理，适合小数据或频繁创建/销毁的场景。
- 直接内存：通过`ByteBuffer.allocateDirect()`分配，避免JVM堆与操作系统间的数据复制，提升I/O性能，但需手动释放。
- 池化（Pooled）与非池化（Unpooled）：
 - 池化内存（`PooledByteBufAllocator`）通过预分配内存块并复用，减少频繁分配/回收的开销。
 - 非池化内存（`UnpooledByteBufAllocator`）每次分配新内存，适合临时使用场景。

2. 内存规格划分

Netty将内存划分为四类以适应不同需求：

- Tiny（<512B）：用于小对象（如协议头）。
- Small（<8KB）：适用于常见数据包。
- Normal（8KB-16MB）：通过`PoolChunk`管理，按页（8KB）分配。
- Huge（>16MB）：直接分配非池化内存，避免大块内存浪费。

内存分配与回收机制

1. 分层内存池结构

- PoolArena：管理内存池，分为`HeapArena`（堆内存）和`DirectArena`（直接内存）。
- PoolChunk：以16MB为单元，拆分为多个8KB的页（Page），按需分配内存块。
- PoolSubpage：将8KB的页进一步拆分为更小的块（如512B），提升小内存利用率。
- PoolThreadCache：线程本地缓存，减少多线程竞争，优先从本地缓存分配内存。

2. 引用计数与释放

- 引用计数（Reference Counting）：每个`ByteBuf`维护计数器，通过`retain()`增加、`release()`减少引用，归零时触发释放。
- 自动回收：池化内存释放后回归内存池；非池化堆内存由GC回收，直接内存需显式释放。
- 泄漏检测：通过`ResourceLeakDetector`监控未释放的`ByteBuf`，支持`PARANOID`模式记录泄漏堆栈。

零拷贝优化

1. 零拷贝技术

- CompositeByteBuf：逻辑合并多个`ByteBuf`，避免物理复制（如合并协议头和数据体）。

- FileRegion：通过 `FileChannel.transferTo()` 直接传输文件，减少内核态与用户态的数据拷贝。
- Direct Buffer 复用：网络传输直接使用堆外内存，省去堆内到堆外的复制步骤。

背景问题解决方案

1. 设置流量控制机制

高低水位 (WriteBufferWaterMark)：

```
channel.config().setWriteBufferHighWaterMark(64 * 1024); // 设置高水位
if (channel.isWritable()) {
    channel.writeAndFlush(data); // 可写时发送
} else {
    // 丢弃或者存储至内存或redis适时重发, 第一阶段建议直接丢弃
}
```

2. 连接channel关闭,对于长时间处于弱网的连接进行关闭

记录channel超过高水位的次数,如果长时间处于弱网,关闭连接.

3. 强制释放策略：

在exceptionCaught和channelInactive事件中释放未发送的积压数据：

```
@Override
public void exceptionCaught(ChannelHandlerContext ctx, Throwable cause) {
    ChannelOutboundBuffer buffer = ctx.channel().unsafe().outboundBuffer();
    buffer.failFlushed(cause, true); // 释放积压数据并关闭连接
}
```